Clark Peng

San Jose, CA | cipeng@ucla.edu | linkedin.com/in/clarkpeng | github.com/clarkipeng | clarkipeng.github.io

Education

University of California, Los Angeles — Los Angeles, CA
B.S. in Computer Science (GPA: 4.0) — *Expected June 2028*

Technical Skills

Languages: Python, C++17/20, SQL, Bash, TypeScript, Mojo, C#
ML Systems: PyTorch, FSDP, FlashAttention, vLLM/SGLang, distributed training, GPU programming (CUDA/Mojo)
Systems / Infra: Linux, GDB, CMake, Docker, PostgreSQL, AWS, Azure

Experience

Catapult (YC S26) — San Francisco, CA
CTO — *Jun 2026 – Present*

- Built the multiplayer runtime, orchestration, and replay system for an LLM strategy-game benchmark; evaluated 19 frontier LLMs across 1K games, producing 257M tokens of gameplay data.

Camfer (YC S24) — San Francisco, CA
Software Engineering Intern — *Aug 2025 – Dec 2025*

- Mid/post-trained Qwen-VL models for instruction following; sequence packing and data preparation drove a 10x throughput gain, with further gains from kernel, quantization, and FSDP tuning.
- Built a sharded synthetic-data pipeline generating 600M+ tokens from 800K CAD parts (60% pre-training, 100% SFT); its Onshape support doubled the addressable user base.
- Built async EC2 rollouts for model error correction; a SolidWorks/Onshape replay layer caught dozens of errors integration tests missed.

UCLA NLP Group — Los Angeles, CA
Student Researcher — *Oct 2024 – Present*

- Co-authored **VideoPhy-2**, a video-generation physical-commonsense benchmark; built its data pipeline, growing it 6x and enabling a 7B model that beat Gemini by 127%. ICLR 2026; Best Paper, ICML 2025 World Models Workshop.
- Applied routing-alignment training to MoE models, improving robustness to new languages with better sample efficiency.

Scale AI — Los Angeles, CA
Technical Intern — *Nov 2024 – Aug 2025*

- Wrote 100+ expert competitive-programming solutions used as training data for LLMs at frontier labs.
- Designed judges and rubrics for RL evaluation of coding agents in deterministic Docker environments.

HMC Music Retrieval Lab — Claremont, CA
Research Intern — *Jun 2025 – Aug 2025*

- Developed self-distillation and DPO adaptation methods for MusicGen steering, then evaluated them across 1M+ audio datapoints.

Projects & Awards

Kaggle Competitions | *Top 1% Competitor, 5+ Gold Notebooks (20k+ Forks)* — 2023 – 2025

- Developed a probability-density regression framework that beat segmentation baselines at event-boundary detection.
- Open-sourced it in the 1,877-team CMI competition; adopted by nearly all top-scoring solutions - Silver Medal, first-author paper.

Graphics Projects | *C++, OpenGL, Unity* — 2022 – Present

- Built a C++/OpenGL 4D game engine, a Unity GPU slime simulation running 1M agents at 120 FPS, and C++ raytracers for image and audio rendering.

Leadership

UCLA ACM AI | *Co-President* — Nov 2025 – Present

- Run ACM AI's 800-member community and 40-officer team across research tracks, ML hackathons, and events drawing thousands of participants; secured club sponsors.

Distinctions

ICPC SoCal Regional Top 10 | USACO Platinum | USAPhO Silver