

Clark Peng

San Jose, CA | cipeng@ucla.edu | [linkedin.com/in/clarkpeng](https://www.linkedin.com/in/clarkpeng) | github.com/clarkipeng | clarkipeng.github.io

EDUCATION

University of California, Los Angeles

2024 – Present

B.S. in Computer Science (GPA: 4.0)

Los Angeles, CA

TECHNICAL SKILLS

Languages: Python, C++17/20, SQL, Bash, TypeScript, Mojo, C#

ML Systems: PyTorch, JAX, FSDP, FlashAttention, vLLM, SGLang, distributed training, GPU programming (CUDA/Mojo)

Systems / Infra: Linux, GDB, CMake, C++ STL, C++ concurrency, Docker, PostgreSQL, AWS, Azure, CoreWeave, Lambda Cloud

EXPERIENCE

Origami (YC S26)

San Francisco, CA

CTO

2026 – Present

- Lead architecture and development of AI research infrastructure for cloud GPU training, experiment tracking, observability, artifact storage, and run lifecycle management.
- Built AWS ECS-backed services to launch and manage GPU runs across Lambda Cloud, AWS EC2, CoreWeave, and Azure, including inspection, cancellation, cleanup, and debugging workflows.

Camfer (YC S24)

San Francisco, CA

Software Engineering Intern

Aug 2025 – Dec 2025

- Scaled multimodal training for Qwen2.5-VL and Qwen3-VL (8B/30B) on 8 H100 GPUs using FSDP and multimodal sequence packing, increasing throughput 10x.
- Built a sharded synthetic-data pipeline that generated 600M+ tokens from 800K CAD parts and reduced preprocessing and tokenization from hours to minutes.
- Built asynchronous EC2 rollout system for model error correction and a 10,000+ line SolidWorks/Onshape replay layer that caught dozens of errors missed by integration tests.

UCLA NLP Group

Los Angeles, CA

Student Researcher

Oct 2024 – Present

- Co-authored **VideoPhy-2** (ICLR 2026; Best Paper, ICML 2025 WorldModel Workshop), a physical-commonsense benchmark for video generation.
- Optimized 30B MoE (3B active) routing-alignment training by packing teacher and student batches, reducing runtime and memory usage while achieving a 2x speedup.

Scale AI

Los Angeles, CA

Technical Intern

Nov 2024 – Aug 2025

- Authored 100+ expert-level solutions to problems rated 2000+ for LLM coding-evaluation datasets.
- Built 10+ deterministic Docker test environments for open-source Python repositories used in RLVR rubrics and autonomous coding-agent evaluation.

PROJECTS & AWARDS

Competitive Programming | *ICPC Top 10, USACO Platinum*

2021 – Present

- Placed Top 10 at the ICPC Southern California Regional; solved timed algorithmic problems across DP, geometry, graph theory, and game theory.

Kaggle Competitions | *Top 1% Competitor, Top 100 Notebooks*

2023 – 2025

- Developed a probability-density regression framework for event-boundary detection; methodology was adopted by top competitors and outperformed standard segmentation approaches on external datasets.
- Published 5+ gold-medal notebooks with 20k+ forks on model training, inference speed, and data augmentation.

Graphics Projects | *C++, OpenGL, Unity*

2022 – Present

- Built a custom C++/OpenGL 4D game engine with BVHs, Verlet sphere physics, and a custom shader; simulated 1M GPU slime agents at 120 FPS; built C++ raytracers for image rendering and audio propagation at 660k rays/min.

LEADERSHIP

UCLA ACM AI | *President*

Nov 2025 – Present

- Lead UCLA's largest AI organization (800+ members), managing 30 student leaders and a \$5k annual budget across research projects, ML hackathons, and collaborations with ICPC, Studio, Hack, and Cyber.